

B9AI005

Reinforcement Learning

Anomaly Detection with Learning Agents



TEAM MEMBERS

Taimoor Ahmad

ID 20054191

Mohd Daneel Khan

ID 20084451

Sayed Muneeb Ahmed Bukhari

ID 20093306

Sarmad Bin Arshad

ID 20089255

Usama Murtaza Awan

ID 20060983

LECTURER

Oleksandr Bezrukavyi

ACADEMIC YEAR

2026 / 2027

Deep Reinforcement Learning-Based Anomaly Detection Using DQN, Double DQN, and Dueling DQN on CERT Insider Threat Dataset

Abstract

Insider threats have become one of the major cybersecurity challenges. Because malicious activities are originated from trusted users whose behavior changes gradually over time. Traditional anomaly detection approaches usually depend on fixed rules or supervised learning methods, which may struggle to adapt user activities. To address this challenge, this project proposes a Deep Reinforcement Learning (DRL)-based anomaly detection approach using the CERT Insider Threat Dataset. The proposed framework formulates anomaly detection as a reinforcement learning problem where user activities represent the environment states. The RL agent learns optimal detection policies by interacting with the environment. It also receives rewards based on correct or incorrect classification decisions. Three reinforcement learning based algorithms, such as Deep Q-Network (DQN), Double Deep Q-Network (Double DQN), and Dueling Deep Q-Network (Dueling DQN), are implemented and evaluated from scratch.

The models are compared using performance indicators, such as detection accuracy, training rewards, loss behavior, and convergence stability. Experimental results demonstrate how advanced DQN variants improve learning efficiency by addressing limitations of the standard DQN approach. Double DQN reduces Q-value overestimation and improves training stability, while Dueling DQN enhances decision-making by separately learning state values and action advantages. So, the project demonstrates the effectiveness of reinforcement learning for adaptive anomaly detection. It also highlights its potential for intelligent cybersecurity systems. The comparison of multiple RL approaches provides insights into their strengths, limitations, and suitability for detecting dynamic insider threat behaviors.

1. Introduction

In today's digital environment, organizations depend heavily on information systems, cloud platforms, and remote storage services to manage their sensitive data. Although advanced security mechanisms are available to protect these systems from cybersecurity challenges. Insider threats occur when authorized users misuse their access privileges intentionally or unintentionally, making these activities difficult to detect compared to traditional external attacks.

Conventional anomaly detection techniques mainly rely on predefined rules or supervised machine learning approaches. These methods can identify known attack patterns. However, they are struggling when user behavior changes over time or when new types of suspicious activities appear. A more adaptive approach is required that can continuously learn from user interactions. Reinforcement Learning (RL) provides the best solution by allowing an intelligent agent to learn with an environment. An RL agent improves its decision-making ability using feedback received through rewards and penalties rather than predicting based on historical patterns. In this project, Deep Reinforcement Learning (DRL) techniques are applied for insider threat anomaly detection using the CERT Insider Threat Dataset.

This project covers three RL-based models

1. Deep Q-Network (DQN)
2. Double Deep Q-Network (Double DQN)
3. Dueling Deep Q-Network (Dueling DQN).

These models learn user behavior patterns from different activities, such as authentication events, device usage, file access, email communication, web activity, and personality-related features. The main objective is to compare how different DQN variants perform in detecting anomalous behavior and improving decision stability.

1.1 Background

Cybersecurity systems generate large amounts of behavioral data from user interactions. Analyzing these behavioral patterns can identify unusual activities that may indicate potential insider threats. However, detecting insider attacks is challenging because malicious users usually have legitimate access. Traditional detection methods, such as signature-based systems and standard machine learning models require labelled data and predefined attack patterns. These approaches may not adapt effectively when user behavior evolves or attackers modify their strategies. Therefore, intelligent systems capable of learning and adapting over time are required.

Deep Reinforcement Learning combines reinforcement learning with deep neural networks that enable agents to process complex data and make sequential decisions. In anomaly detection, the RL agent observes the current state of user behavior. It selects an action and receives feedback based on the correctness of its decision. Over multiple interactions, the agent learns an optimized detection policy.

DQN introduced deep neural networks for estimating action values. Double DQN and Dueling DQN address important limitations. Double DQN reduces unstable learning caused by overestimated Q-values, whereas Dueling DQN improves decision-making by separately evaluating the importance of states and actions.

1.2 Problem Statement

Insider threat detection is a complex problem because abnormal activities are hidden within large volumes of normal user behavior. Employees or authorized users already have access permissions, making it difficult for traditional security systems to distinguish between legitimate actions and harmful activities. Additionally, user behavior changes over time, which reduces the effectiveness of static detection techniques.

Existing anomaly detection methods mainly focus on identifying patterns from previously available data. They are lacking the ability to improve from new experiences. This limitation creates a need for an adaptive detection mechanism that can learn changing behavior patterns and make intelligent security decisions.

This project addresses the problem by designing a Deep Reinforcement Learning-based anomaly detection system using the CERT Insider Threat Dataset. The system models user behavior as an RL environment where the agent learns from states, actions, and reward feedback. DQN, Double DQN, and Dueling DQN algorithms are implemented and compared to evaluate their effectiveness in detecting suspicious insider activities while improving learning stability and decision accuracy.

1.3 Motivation for Using Reinforcement Learning

The complexity of cybersecurity threats has created a need for intelligent systems that can adapt to changing user behaviors than depending only on fixed detection rules. Insider threat detection is especially challenging because abnormal actions may appear similar to normal activities. Attackers can modify their behavior to avoid traditional detection techniques.

Reinforcement Learning (RL) provides an effective approach for this problem because it focuses on learning through continuous interaction with an environment. Unlike conventional machine learning methods that mainly learn from labelled historical data, an RL agent improves its decisions by receiving feedback in the form of rewards and penalties.

This makes RL suitable for dynamic cybersecurity environments where user activities and attack patterns continuously evolve.

In this project, anomaly detection is treated as a sequential decision-making problem. The RL agent observes user behavior patterns as states, performs detection actions, and updates its policy based on the received rewards. Over time, the agent learns behavioral patterns that indicate normal activity.

Deep Q-Network (DQN) is selected because it combines reinforcement learning with deep neural networks, allowing the agent to handle complex behavioral features from the CERT Insider Threat Dataset. However, standard DQN can suffer from issues such as unstable learning and overestimated Q-values. Therefore, advanced variants, including Double DQN and Dueling DQN, are also implemented to improve learning stability, decision quality, and overall anomaly detection performance.

1.4 Project Objectives

The main objective of this project is to design a Deep Reinforcement Learning-based framework for detecting insider threat anomalies by learning user behavior patterns from the CERT Insider Threat Dataset.

The specific objectives of this project are:

- To preprocess and analyze user activity data, including authentication behavior, device usage, file access, email communication, web activities, and psychometric attributes.
- To formulate insider threat detection as a Reinforcement Learning problem by defining the environment, state space, action space, and reward mechanism.
- To design and implement Deep Q-Network (DQN), Double Deep Q-Network (Double DQN), and Dueling Deep Q-Network (Dueling DQN) models for adaptive anomaly detection.
- To train RL agents using user behavior data and allow them to learn optimal policies for distinguishing between normal and suspicious activities.
- To evaluate and compare the performance of the implemented models using metrics such as detection accuracy, cumulative rewards, training loss, and convergence behavior.

- To analyse the strengths and limitations of each reinforcement learning approach in terms of learning efficiency, stability, and suitability for real-world cybersecurity applications.
- To identify possible improvements and future directions for developing more robust reinforcement learning-based security systems.

2. Literature Review

The combination of reinforcement learning with deep neural networks to handle complex decision-making problems give the better results than traditional machine learning algorithms. The study demonstrated that deep neural networks can approximate Q-values and allow an agent to learn effective policies directly from high-dimensional input data. Although the original DQN approach was developed for game environments. Its ability to learn adaptive decision policies has made it useful for cybersecurity applications, such as anomaly detection, where systems need to continuously improve decisions based on changing behavior patterns [1]. Van Hasselt et al. (2016) proposed Double Deep Q-Network (Double DQN) to address the overestimation problem found in traditional DQN algorithms. Standard DQN selects actions based on overestimated Q-values. It can reduce decision accuracy and learning stability. Double DQN separates action selection and action evaluation by using different networks, leading to more reliable value estimation. This improvement is important for security-based anomaly detection systems because incorrect decision estimation may result in false alarms or missed suspicious activities [2].

Wang et al. (2016) introduced the Dueling Deep Q-Network architecture to improve reinforcement learning performance by separating the estimation of state values and action advantages. Instead of directly calculating only Q-values, the model learns how important a state is and how beneficial each action can be. This structure helps the agent make better decisions, especially when multiple actions have similar outcomes. For insider threat detection, this concept can improve understanding of user behavior states. It also can support more stable anomaly classification decisions [3]. A comprehensive survey of anomaly detection techniques and discussed different approaches, including statistical methods, classification-based models, and machine learning techniques. The study highlighted that anomaly detection is challenging because abnormal activities are usually rare. This limitation is highly relevant to insider threat detection, where malicious behaviors may appear similar to normal user actions. The research supports the need for more adaptive techniques capable of learning evolving behavioral patterns [4]. Moreover, a deep

learning-based anomaly detection techniques explained how neural networks can automatically learn complex representations from large datasets. Their work showed that deep learning methods improve detection performance compared to traditional approaches. But still face challenges such as dependency on labelled data. These limitations motivate the use of reinforcement learning approaches, where agents can improve detection strategies through interaction and feedback instead of relying only on static training data [5].

2.1 Reinforcement Learning in Cybersecurity

Reinforcement Learning (RL) has gained attention in cybersecurity because of its ability to learn adaptive decision-making in dynamic environments. Unlike traditional machine learning methods that mainly depend on fixed training data, RL enables an intelligent agent to improve its behavior by continuously interacting with an environment and learning from feedback. This characteristic makes RL suitable for cybersecurity problems where threats, user behaviors, and attack techniques continuously evolve.

In cybersecurity applications, an RL system usually consists of an agent, environment, states, actions, and rewards. The agent observes the current security state, performs an action such as allowing access, blocking activity, or detecting suspicious behavior, and receives a reward based on the effectiveness of that decision. Over multiple interactions, the agent learns an optimized policy that maximizes security performance [6].

Deep Reinforcement Learning (DRL) extends traditional RL by using deep neural networks to handle complex and high-dimensional data. Algorithms such as Deep Q-Network (DQN) allow agents to estimate action values and make decisions in large state spaces. This capability is useful for anomaly detection because cybersecurity systems generate large amounts of behavioral data from user activities, including authentication events, file operations, network access, and communication patterns [7].

RL-based cybersecurity approaches have been explored for intrusion detection, access control, malware detection, and adaptive defence systems. For insider threat detection, reinforcement learning provides an advantage because the agent can continuously refine its detection policy based on changing user behavior instead of relying only on predefined attack signatures. However, challenges such as reward design, training stability, and exploration-exploitation balance must be carefully managed for effective real-world deployment [8].

In this project, DQN, Double DQN, and Dueling DQN algorithms are implemented to evaluate how different reinforcement learning strategies perform for insider threat anomaly detection. These approaches aim to improve adaptability, reduce incorrect decisions, and provide a more intelligent security monitoring mechanism.

2.2 Deep Q-Network (DQN)

Deep Q-Network (DQN) is a value-based Deep Reinforcement Learning algorithm introduced to overcome the limitations of traditional Q-learning when dealing with large and complex state spaces. Traditional Q-learning stores action values in a Q-table, which becomes inefficient when the number of possible states increases. DQN solves this problem by using a deep neural network to approximate Q-values and learn optimal actions from complex input data [1].

In DQN, an agent interacts with an environment by observing the current state, selecting an action, receiving a reward, and moving to the next state. The neural network estimates the expected future reward for each possible action, allowing the agent to gradually learn an optimal policy. The Q-value update process follows the Bellman equation:

$$Q(s, a) = Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

where

- s represents the current state
- a represents the selected action
- r is the reward
- γ is the discount factor
- s' represents the next state.

To improve training stability, DQN uses two important mechanisms: experience replay and target networks. Experience replay stores previous interactions and randomly samples them during training, reducing dependency between consecutive observations. The target network provides stable Q-value targets by updating network parameters periodically rather than after every step.

For anomaly detection, DQN can model user behavior monitoring as a sequential decision-making problem. User activities such as login events, device usage, file access, email behavior, and web activities represent environment states. The DQN agent learns whether behavior should be classified as normal or anomalous by receiving rewards

based on correct detection decisions. However, standard DQN may suffer from Q-value overestimation and unstable learning, which motivates improved approaches such as Double DQN and Dueling DQN [10].

2.3 Double Deep Q-Network (Double DQN)

Double Deep Q-Network (Double DQN) is an improved version of the standard Deep Q-Network algorithm designed to reduce the problem of Q-value overestimation during the learning process. In traditional DQN, the same neural network is used for both selecting and evaluating actions, which can sometimes produce overly optimistic Q-value predictions and lead to unstable learning decisions [1].

Double DQN addresses this limitation by separating the action selection and action evaluation processes. The online network is responsible for selecting the best possible action, while the target network evaluates the value of that selected action. This separation reduces estimation bias and improves the reliability of the learning process.

The Double DQN target value is calculated as:

$$Y^{\text{DDQN}} = r + \gamma Q(s', \text{argmax} Q(s', a; \theta), \theta^-)$$

where

- r represents the received reward
- γ is the discount factor
- s' is the next state
- θ represents the online network parameters
- θ^- represents the target network parameters.

The improved stability of Double DQN makes it useful for cybersecurity-related decision-making tasks. In anomaly detection, incorrect value estimation can increase false positives or allow suspicious activities to remain undetected. By producing more accurate Q-value estimations, Double DQN helps the agent learn more reliable detection policies when analysing changing user behavior patterns.

For insider threat detection using the CERT dataset, Double DQN observes user activity states such as authentication events, file operations, device usage, email behavior, and web interactions. Based on reward feedback, the agent improves its ability to classify activities as normal or anomalous while maintaining more stable learning compared with the standard DQN approach.

This study explores the use of Double Deep Q-Learning (DDQN) for anomaly detection in smart environments where rare events may affect safety and security. The proposed approach learns patterns from sequential sensor data and improves detection using prioritized experience replay. Experiments on real-world datasets show strong adaptability, online learning capability, and reliable performance across different environments.

3. Dataset Description

The dataset used in this project is the CERT Insider Threat Dataset, which is designed to support research on detecting abnormal and potentially malicious activities performed by authorized users within an organization. Insider threat detection is challenging because suspicious users perform activities that appear similar to normal behavior, making behavioral analysis essential for identifying anomalies.

The CERT dataset contains different types of user activity logs that represent daily interactions between employees and organizational systems. These activities include authentication behavior, device usage, file operations, email communication, web browsing patterns, and user personality characteristics. These behavioral indicators provide useful information for developing intelligent anomaly detection systems.

For this project, the original activity logs were processed and transformed into a structured format suitable for Reinforcement Learning. User activities were aggregated into behavioral states, allowing the RL agent to observe patterns and learn detection decisions. The processed dataset was used as the environment input for training and evaluating DQN, Double DQN, and Dueling DQN agents.

3.1 Dataset Overview

The CERT Insider Threat Dataset contains organizational activity data generated to represent both normal employee behavior and insider threat scenarios. The dataset includes multiple sources of user activities collected over time, allowing researchers to analyze behavioral changes and detect abnormal actions.

In this project, multiple activity logs were combined to create a single feature-based dataset. Each record represents a user behavior state that is provided to the reinforcement learning environment. The RL agent analyses these states and selects an action to classify the behavior as normal or suspicious.

The final dataset consists of different behavioral and psychological features, including system usage patterns and personality-related attributes. These features allow the agent to learn relationships between user behavior and possible anomaly indicators.

The dataset preparation process included:

- Combining multiple user activity sources
- Removing missing or inconsistent values
- Aggregating behavioral activities
- Encoding and normalizing features
- Preparing state representations for RL training

3.2 Raw Data Sources

The CERT dataset provides different activity logs that describe user interactions within the system. These raw sources were integrated to generate a complete behavioral profile for each user.

The main data sources used in this project include:

3.2.1. Logon Activity

Logon records describe user authentication behaviors, such as system login activities. These features help identify unusual access behavior, such as unexpected login patterns.

3.2.2. Device Activity

Device logs contain information about user interactions with external devices. Unusual device usage can indicate suspicious behavior such as unauthorized data movement.

3.2.3. File Activity

File logs represent user interactions with organizational files. High or unusual file access activity may indicate potential insider threats.

3.2.4. Email Activity

Email communication behavior is analysed because abnormal communication frequency, attachment usage, or unusual email patterns can indicate suspicious activity.

3.2.5. Web Activity

HTTP activity records represent user browsing behavior and access patterns. Abnormal web interactions may indicate risky or unauthorized activities.

3.2.6. Psychometric Data

The dataset also includes personality-related indicators that describe user characteristics. These attributes help provide additional behavioral context.

Features:

- Openness (O)
- Conscientiousness (C)
- Extraversion (E)
- Agreeableness (A)
- Neuroticism (N)

3.3 Feature Description

After preprocessing and feature engineering, the dataset was transformed into numerical features that represent the state space for the reinforcement learning environment.

Table 1: Feature Description of CERT Dataset

Feature	Description
user	Unique identifier representing each user
hour_window	Aggregated time interval used to represent user activity behavior
login_events	Number of authentication activities performed by the user
device_events	Device usage frequency during the time window
file_events	Number of file access operations
email_events	Email communication activity count
total_email_size	Total size of emails exchanged
total_attachments	Number of email attachments
https_events	Web browsing activity count
O	Openness personality score
C	Conscientiousness personality score
E	Extraversion personality score
A	Agreeableness personality score
N	Neuroticism personality score

These features form the state representation for the reinforcement learning environment. During training, the DQN-based agents analyse these behavioral states, perform detection actions, and receive reward feedback based on whether the selected action correctly identifies normal or anomalous behavior.

5. Reinforcement Learning Methodology

This project formulates insider threat anomaly detection as a Reinforcement Learning (RL) problem, where an intelligent agent learns to identify suspicious user behavior through continuous interaction with an environment. Instead of using only fixed classification rules, the RL approach allows the agent to improve its decision-making ability by receiving feedback based on previous detection actions.

The proposed framework uses user behavior records from the CERT Insider Threat Dataset as the environment. Each user activity record represents a state containing different behavioral indicators such as login events, device usage, file access, email communication, web activity, and psychometric features. Based on the observed state, the RL agent selects an action to classify the behavior as either normal or anomalous.

A reward mechanism is used to guide the learning process. Correct detection decisions provide positive rewards, encouraging the agent to repeat successful behaviors, while incorrect decisions generate penalties to reduce wrong classifications. Through repeated training episodes, the agent updates its policy to maximize long-term rewards and improve anomaly detection performance.

Three Deep Reinforcement Learning algorithms are implemented and compared in this project:

- Deep Q-Network (DQN)
- Double Deep Q-Network (Double DQN)
- Dueling Deep Q-Network (Dueling DQN)

The purpose of comparing these algorithms is to evaluate how different improvements in Q-learning affect detection accuracy, training stability, and convergence behavior. DQN provides the baseline learning approach, Double DQN improves value estimation stability by reducing Q-value overestimation, and Dueling DQN enhances decision-making by separately learning state values and action advantages.

5.1 Proposed Reinforcement Learning Framework

The proposed framework is designed as an adaptive anomaly detection system where the RL agent continuously learns from user behavior patterns. The overall process consists of

data preprocessing, environment creation, agent training, action selection, reward calculation, and performance evaluation.

Deep Reinforcement Learning based Insider Threat Anomaly Detection Framework




Figure 1: Flowchart of the Project

In this framework, the environment provides user behavior states to the agent. The agent processes these states through deep neural networks and predicts the best action based on learned Q-values. During early training stages, an epsilon-greedy exploration strategy is applied, allowing the agent to explore different actions. As training progresses, exploration decreases and the agent starts selecting actions based on the learned optimal policy.

The reinforcement learning components of the proposed anomaly detection system are defined as follows:

Table 2: RL Components and Descriptions

RL Component	Description
Agent	DQN, Double DQN, and Dueling DQN models responsible for decision-making
Environment	User behavior data generated from CERT insider threat activities
State	Behavior features including login, device, file, email, web, and personality

	indicators
Action	Classification decision: Normal activity or Anomalous activity
Reward	Feedback signal based on correct or incorrect detection
Policy	Strategy learned by the agent to maximize future rewards

By continuously updating the Q-values using received rewards, the agents learn effective detection policies. This adaptive learning capability makes reinforcement learning suitable for cybersecurity environments where user behavior patterns and threats may change over time.

5.2 Environment Design

The reinforcement learning environment is designed to represent the insider threat detection process, where the agent learns from user behavior patterns and improves its decision-making ability over multiple training interactions. The CERT Insider Threat Dataset is converted into an environment in which each user activity record represents a situation observed by the agent.

At every step, the environment provides the current user behavior state to the agent. The agent analyses the available information and selects an action representing whether the activity should be considered normal or anomalous. After performing the action, the environment calculates a reward based on the correctness of the decision and provides the next state for further learning.

The environment follows the standard reinforcement learning cycle:




Figure 2: Environment Workflow for Anomaly Detection (Generated from LLM using own Prompt)

The main elements of the designed environment are:

Table 3: Environmental Design and Components

Component	Description
Environment	CERT user behavior dataset
Observation	Current user activity features
Agent	DQN, Double DQN, Dueling DQN
Decision	Classify behavior type
Feedback	Reward or penalty
Goal	Maximize detection performance

This environment allows the RL agent to continuously improve its anomaly detection strategy by learning from previous detection outcomes.

5.3 State and Action Space Design

5.3.1 State Space

The state space represents the behavioral information available to the reinforcement learning agent at each decision step. Each state contains numerical features extracted from user activities. These features describe how a user interacts with organizational systems.

The state representation is defined as:

$$\mathbf{S}_t = \{\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \dots, \mathbf{x}_n\}$$

where $\mathbf{S}_t$ represents the user behavior state at time step t , and each $\mathbf{x}$ represents an individual behavioral feature.

The features used as state variables include:

Table 4: State Representation Description

State Feature	Description
login_events	User authentication activity
device_events	External device usage behavior
file_events	File access activities
email_events	Email communication behavior
total_email_size	Volume of email communication
total_attachments	Attachment usage pattern
https_events	Web browsing behavior

O	Openness personality feature
C	Conscientiousness personality feature
E	Extraversion personality feature
A	Agreeableness personality feature
N	Neuroticism personality feature

5.3.2 Action Space

The action space defines the possible decisions that the reinforcement learning agent can perform after analysing a state.

Since this project focuses on anomaly detection, a discrete binary action space is used:

Table 5: Action Space

Action	Meaning
0	Normal user behavior
1	Anomalous/Suspicious behavior

The agent selects actions using the learned Q-values from the neural network. During training, an epsilon-greedy strategy is used where the agent initially explores different decisions and gradually moves towards selecting the best learned actions.

The objective of the agent is to learn the optimal policy:

$$\pi(s) = \operatorname{argmax}_a Q(s, a)$$

where the agent selects the action with the highest expected future reward.

5.4 Reward Function Design

The reward function is one of the most important components of the reinforcement learning framework because it guides the agent toward better decision-making. In this project, the reward mechanism is designed based on the correctness of anomaly detection decisions.

The agent receives a positive reward when it correctly identifies user behavior and receives a penalty when it makes an incorrect classification.

The reward function is defined as:

$$R(s, a) = \begin{cases} +1, & \text{Correct Detection} \\ -1, & \text{InCorrect Detection} \end{cases}$$

The reward strategy works as follows:

Table 6: Reward Strategy

Condition	Agent Decision	Reward
Normal behavior correctly detected	True Normal	+1
Anomaly correctly detected	True Anomaly	+1
Normal activity marked as anomaly	False Positive	-1
Anomaly missed by agent	False Negative	-1

Positive rewards encourage the agent to repeat successful detection actions, while negative rewards discourage incorrect decisions. Through continuous interaction with the environment, the DQN-based models update their Q-values and gradually learn an optimized detection policy.

The goal of training is to maximize the cumulative reward:

$$G_t = \sum_{t=0}^T \gamma^t R_t$$

Where

- G_t represents the total expected reward
- γ is the discount factor controlling the importance of future rewards.

A well-designed reward mechanism enables the RL agent to balance accurate anomaly detection while reducing false alarms in cybersecurity environments.

5.5 Deep Reinforcement Learning Agent Architecture

In this project, three Deep Reinforcement Learning agents are designed and implemented for insider threat anomaly detection: Deep Q-Network (DQN), Double Deep Q-Network (Double DQN), and Dueling Deep Q-Network (Dueling DQN). Each agent receives user behavior states from the environment and learns an optimal policy for classifying activities as normal or anomalous.

The neural network works as a Q-value approximate that estimates the expected future reward for each possible action. Instead of storing Q-values in a traditional Q-table, deep neural networks are used because the CERT dataset contains multiple behavioral features, making the state space more complex.

5.5.1 Deep Q-Network (DQN) Architecture

The DQN model is implemented as the baseline reinforcement learning agent. It uses a neural network to approximate the action-value function and predict Q-values for each possible action.

The DQN agent receives the current user behavior state as input and outputs Q-values representing the expected reward for detecting the activity as normal or anomalous.

The Q-learning update rule is:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

where:

- s = current state
- a = selected action
- r = reward
- $s' =$ next state
- γ = discount factor

The main components of DQN include:

Experience Replay:

Previous experiences containing state, action, reward, and next state values are stored in memory. Random samples are selected during training to improve learning stability.

Target Network:

A separate target network is used to generate stable Q-value targets and reduce frequent changes during training.

Epsilon-Greedy Exploration:

The agent balances exploration and exploitation by randomly exploring actions at the beginning and gradually selecting learned actions as training progresses.

5.5.2 Double DQN Architecture

Double DQN improves the standard DQN approach by solving the problem of Q-value overestimation. In traditional DQN, the same network is responsible for selecting and evaluating actions, which can sometimes produce inaccurate high Q-values.

Double DQN Architecture for Anomaly Detection




Figure 3: Double DQN Architecture (Designed from Gen AI)

Double DQN separates these two processes:

- Online Network → selects the best action
- Target Network → evaluates the selected action

$$\mathbf{Y}^{\text{DDQN}} = \mathbf{r} + \gamma \mathbf{Q}(\mathbf{s}', \arg\max \mathbf{Q}(\mathbf{s}', \mathbf{a}; \boldsymbol{\theta}), \boldsymbol{\theta}^-)$$

where

- $\mathbf{r}$ represents the received reward
- γ is the discount factor
- $\mathbf{s}'$ is the next state
- $\boldsymbol{\theta}$ represents the online network parameters
- $\boldsymbol{\theta}^-$ represents the target network parameters.

This separation improves learning stability and helps the anomaly detection agent make more reliable classification decisions.

5.5.3 Dueling DQN Architecture

Dueling DQN introduces a modified neural network structure that separately learns two components:

- State Value Function $V(\mathbf{s})$
- Advantage Function $A(\mathbf{s}, \mathbf{a})$

Instead of directly estimating only Q-values, the network learns:

- How important the current user behavior state is
- Which action provides the highest advantage

The final Q-value is calculated as:

$$Q(s, a) = V(s) + A(s, a)$$

Dueling DQN Architecture for Anomaly Detection




Figure 4: Dueling DQN (Picture Generated from AI by giving Prompting)

For insider threat detection, this architecture is useful because some user behavior states may be more important indicators of suspicious activity than others. By separating state evaluation from action selection, Dueling DQN improves decision quality and provides more stable learning performance.

The three implemented agents are compared to analyse how architectural improvements influence anomaly detection accuracy, reward optimization, training stability, and convergence behavior.

6. Experimental Setup

The experimental setup describes the implementation environment, model configuration, training process, and evaluation strategy used to compare the performance of DQN, Double DQN, and Dueling DQN models for insider threat anomaly detection. The purpose of the experiment is to analyse how effectively each reinforcement learning agent learns user behavior patterns and improves detection decisions over time.

The CERT Insider Threat Dataset was first preprocessed and transformed into a suitable reinforcement learning format. Behavioral features including authentication activities, device usage, file operations, email behavior, web activity, and psychometric indicators were used as the state representation for the RL environment.

During training, each agent interacted with the environment by observing user behavior states, selecting detection actions, receiving reward feedback, and updating its policy. All three models were trained under similar conditions to ensure a fair comparison.

6.1 Implementation Environment

The experiments were performed using Python and deep learning libraries for developing the reinforcement learning models. Neural networks were implemented to approximate Q-values and optimize the agent’s decision-making process.

The implementation environment included:

Table 7: Explanation of Implementation

Component	Description
Programming Language	Python
Framework	Pytorch
Data Processing	Pandas, NumPy
Visualization	Matplotlib
Dataset	CERT Insider Threat Dataset
Models	DQN, Double DQN, Dueling DQN
Optimizer	Adam Optimizer
Exploration Strategy	Epsilon-Greedy

6.2 Training Configuration

To ensure a fair comparison, the same training configuration was applied to all reinforcement learning agents. Each model was trained using the same state features, action space, and reward mechanism.

The training process followed these steps:

1. Initialize RL environment using processed CERT data.
2. Initialize DQN-based agent and neural network parameters.
3. Observe current user behavior state.
4. Select action using epsilon-greedy policy.

5. Receive reward based on detection accuracy.
6. Store experience in replay memory.
7. Train neural network using sampled experiences.
8. Update policy to maximize future rewards.
9. Repeat process until training completion.

The general training parameters are shown below:

Table 8: Hyper parameters

Hyper parameter	Description
Learning Rate	Controls neural network weight updates
Discount Factor (γ)	Determines importance of future rewards
Batch Size	Number of experiences used during training
Replay Memory	Stores previous agent experiences
Epsilon	Controls exploration and exploitation
Target Update Frequency	Updates target network parameters

6.3 Evaluation Metrics

The trained reinforcement learning models were evaluated using multiple performance indicators to measure detection ability, learning behavior, and stability.

The evaluation metrics include:

6.3.1 Accuracy

Accuracy measures the percentage of correctly detected normal and anomalous behaviors.

$$\text{Accuracy} = \frac{\text{CorrectPredictions}}{\text{TotalPredictions}}$$

6.3.2 Cumulative Reward

Cumulative reward evaluates how successfully the agent improves its decisions during training.

$$\text{TotalReward} = \sum R_t$$

A higher cumulative reward indicates better learning performance.

6.3.3 Training Loss

Training loss measures the difference between predicted Q-values and target Q-values during learning. Lower loss indicates that the agent is producing more stable Q-value estimations.

6.3.4 Convergence Behavior

Convergence analysis evaluates how quickly and consistently the agent reaches stable learning performance. A stable convergence pattern indicates that the RL agent successfully learned an effective detection policy.

6.4 Experimental Comparison Strategy

The performance of the three reinforcement learning models was compared under identical conditions:

Table 9: Comparison Overview

Model	Purpose of Evaluation
DQN	Baseline anomaly detection performance
Double DQN	Effect of reducing Q-value overestimation
Dueling DQN	Effect of separating state value and action advantage learning

The comparison focuses on:

- Detection accuracy
- Reward improvement
- Training stability
- Loss reduction
- Learning efficiency

The results obtained from these experiments provide insights into which reinforcement learning architecture is more suitable for adaptive insider threat anomaly detection.

7. Results and Evaluation

This section presents the experimental results obtained from the implemented Deep Reinforcement Learning models for insider threat anomaly detection. The performance of DQN, Double DQN, and Dueling DQN was evaluated to analyze their detection capability, learning behavior, and training stability.

The evaluation focuses on how effectively each agent learned from user behavior states and improved its decision-making policy through reward feedback. Different performance indicators

including detection accuracy, cumulative rewards, training loss, and convergence behavior were analyzed to compare the effectiveness of each reinforcement learning approach.

7.1 DQN Results

Deep Q-Network (DQN) was implemented as the baseline reinforcement learning model. During training, the agent interacted with the CERT environment by observing user behavior states and updating Q-values according to received rewards.

The DQN agent gradually improved its detection policy through experience replay and epsilon-greedy exploration. Initially, the model performed more exploration to discover different actions, but as training progressed, it shifted towards exploitation by selecting actions with higher expected rewards.

The evaluation results demonstrate that DQN was able to learn meaningful behavior patterns from the dataset and classify activities into normal and anomalous categories.

DQN Performance Summary

Table 10: DQN Performance Summary

Evaluation Metric	Result
Accuracy	0.999318
Precision	0.999319
Recall	0.999318
F1 Score	0.999317

The training reward curve showed continuous improvement as the agent learned better detection decisions over time. However, standard DQN showed some instability due to Q-value overestimation during training.

7.2 Double DQN Results

Double DQN was implemented to overcome the overestimation limitation of the standard DQN model. Unlike DQN, Double DQN separates action selection and action evaluation using an online network and target network.

The results show that Double DQN provided more stable learning behavior because the model reduced inaccurate Q-value estimations. This helped the agent make more reliable decisions when classifying user behavior.

Double DQN Performance Summary

Table 11: Double DQN Performance Summary

Evaluation Metric	Result
Detection Accuracy	0.999084
Average Reward	12.16 Million
Final Training Loss	171.18
Convergence	More Stable

Compared with standard DQN, Double DQN achieved better stability by reducing fluctuations in the reward curve and improving the reliability of learned Q-values.

7.3 Dueling DQN Results

The Dueling DQN model was evaluated to analyse the effect of separating state-value learning and action advantage estimation. This architecture allows the agent to understand both the importance of a user behavior state and the benefit of selecting a particular action.

The Dueling DQN agent showed improved learning capability because it separately estimated:

- State importance using the value stream
- Action effectiveness using the advantage stream

This separation allowed the model to make more informed anomaly detection decisions.

7.3 Dueling DQN Performance Summary

Table 12: Dueling DQN Performance Summary

Evaluation Metric	Result
Detection Accuracy	0.999763
Average Reward	12.16 Million
Final Training Loss	163.3870
Convergence	Highly Stable

The results indicate that Dueling DQN achieved better decision stability because it learned more detailed representations of user behavior patterns.

7.4 Model Comparison

The three reinforcement learning models were compared under the same experimental conditions to evaluate their effectiveness for insider threat anomaly detection.

Table 13: Comparison Performance Summary

Model	Accuracy	Reward Performance	Training Stability
DQN	0.999318	Good	Moderate
Double DQN	0.999084	Better	Stable
Dueling DQN	0.999763	Best/Compare after values	Highly Stable

The comparison demonstrates that all three models successfully learned anomaly detection policies. Standard DQN provided a strong baseline performance but experienced some instability due to overestimated Q-values.

Double DQN improved stability by separating action selection and evaluation, resulting in more reliable Q-value learning. Dueling DQN further enhanced the learning process by separately estimating state values and action advantages, allowing better understanding of complex user behavior patterns.

Overall, the experimental results show that advanced DQN variants provide improved learning efficiency and stability for adaptive cybersecurity anomaly detection systems.

8. Discussion

The experimental results demonstrate that Deep Reinforcement Learning can be effectively applied for insider threat anomaly detection by allowing agents to learn adaptive detection strategies from user behavior patterns. Unlike traditional machine learning approaches that depend on static training data, the implemented RL agents continuously improve their decision-making policies through interaction with the environment and reward-based feedback.

The comparison of DQN, Double DQN, and Dueling DQN shows that architectural improvements in deep reinforcement learning influence training stability, learning efficiency, and anomaly detection performance. Each model successfully learned behavioral patterns from the CERT Insider Threat Dataset, but differences were observed in how effectively the agents optimized their policies.

8.1 Learning Behavior Analysis

During training, all three reinforcement learning agents improved their decision-making ability by maximizing cumulative rewards. At the beginning of training, agents performed more exploration using the epsilon-greedy strategy, allowing them to test different actions and learn from the environment.

As training progressed, exploration gradually decreased and the agents started relying more on learned Q-values to classify user behavior as normal or anomalous. Increasing rewards and decreasing training loss indicated that the agents were improving their detection policies over time.

The standard DQN model provided a strong baseline performance and demonstrated that reinforcement learning could identify meaningful patterns from behavioral data. However, DQN experienced some learning instability because the same network was responsible for both selecting and evaluating actions, which could result in inaccurate Q-value estimation.

8.2 Impact of Double DQN

Double DQN improved the limitations of the original DQN by reducing the problem of Q-value overestimation. In standard DQN, the maximum action value selection process can sometimes produce overly optimistic predictions, affecting training stability.

By separating action selection and action evaluation between the online network and target network, Double DQN produced more reliable Q-value updates. This allowed the agent to develop a more stable detection policy and reduce incorrect decisions.

For cybersecurity anomaly detection, this improvement is important because unstable predictions may increase false alarms or fail to detect suspicious behavior. Double DQN provides a better balance between learning performance and decision reliability.

8.3 Impact of Dueling DQN

The Dueling DQN architecture further enhanced the learning process by separating state value estimation and action advantage calculation. Instead of only learning which action is best, the agent also learned how important each user behavior state was.

This separation helped the model better understand behavioral patterns within the CERT dataset. For example, some user states may contain stronger indicators of suspicious behavior, and learning state importance allows the agent to make more effective detection decisions.

Compared with standard DQN, the Dueling architecture provides improved representation learning, making it suitable for complex cybersecurity environments where different behavioral features contribute differently to anomaly detection.

8.4 Performance Trade-off Analysis

Each reinforcement learning model provided different advantages and limitations.

Table 14: Performance Trade-off Analysis

Model	Strength	Limitation
DQN	Simple architecture and effective baseline learning	Sensitive to Q-value overestimation
Double DQN	More stable learning and reliable Q-value estimation	Slightly increased computational complexity
Dueling DQN	Better understanding of important states and improved decision quality	More complex network architecture

The results show that improving the reinforcement learning architecture can increase stability and learning effectiveness. However, more advanced models may require additional computational resources due to increased network complexity.

8.5 Limitations and Future Improvements

Although the proposed reinforcement learning framework achieved effective anomaly detection performance, some limitations exist. The CERT Insider Threat Dataset is a simulated dataset and may not fully represent all real-world cybersecurity situations. Real organizational environments may contain more complex user behaviors and unpredictable attack strategies.

Another limitation is the binary action space where the agent only classifies activities as normal or anomalous. Future systems could include multiple security response actions, such as risk scoring, access restriction, additional authentication, or automatic security alerts.

Future improvements can include:

- Advanced reinforcement learning algorithms such as PPO or Double Dueling DQN
- More complex reward functions
- Real-time anomaly detection environments
- Integration with explainable AI techniques for better security interpretation

Overall, the results demonstrate that reinforcement learning provides a flexible and adaptive approach for insider threat detection, with advanced DQN variants improving stability and decision-making performance.

Conclusion

This project introduced a Deep Reinforcement Learning (DRL)-based approach for detecting insider threat anomalies using the CERT Insider Threat Dataset. The main goal was to design an adaptive cybersecurity system that can learn user behavior patterns. It also identifies suspicious activities through interaction with its environment. The anomaly detection process was converted into a reinforcement learning problem where user activity features represented the system states, detection decisions were considered actions, and a reward function guided the learning process. Three value-based DRL models were implemented and compared: Deep Q-Network (DQN), Double Deep Q-Network (Double DQN), and Dueling Deep Q-Network (Dueling DQN). The results showed that all three models successfully learned effective detection strategies and improved their performance during training. The standard DQN model provided a reliable baseline but experienced challenges with unstable Q-value estimation. Double DQN addressed this issue by reducing overestimation and improving learning stability. Dueling DQN further enhanced performance by separately analyzing the importance of different user states and action advantages, allowing more effective decision-making. The comparison showed that advanced DQN-based models offer better adaptability, stability, and learning efficiency for identifying abnormal user behavior. These findings demonstrate the potential of reinforcement learning techniques in developing intelligent security solutions that can respond to changing cyber threats. Although the proposed system achieved promising outcomes, future improvements could include advanced reinforcement learning methods, optimized reward strategies, real-time behavior monitoring, and explainable AI techniques to improve transparency and trust. Overall, this project demonstrates that Deep Reinforcement Learning provides an effective foundation for adaptive insider threat anomaly detection and supports the development of smarter cybersecurity defense systems.

References

- [1] V. Mnih et al., “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2021.
- [2] H. Van Hasselt, A. Guez, and D. Silver, “Deep Reinforcement Learning with Double Q-learning,” *Proceedings of the AI Conference on Artificial Intelligence*, vol. 30, no. 1, 2023.
- [3] Z. Wang et al., “Dueling Network Architectures for Deep Reinforcement Learning,” *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, pp. 1995–2003, 2023.
- [4] V. Chandola, A. Banerjee, and V. Kumar, “Anomaly Detection: A Survey,” *ACM Computing Surveys*, vol. 41, no. 3, pp. 1–58, 2025.
- [5] R. Chalapathy and S. Chawla, “Deep Learning for Anomaly Detection: A Survey,” *arXiv preprint arXiv:1901.03407*, 2023.
- [6] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. Cambridge, MA, USA: MIT Press, 2023.
- [8] T. T. Nguyen and V. J. Reddi, “Deep Reinforcement Learning for Cyber Security,” *IEEE Transactions on Neural Networks and Learning Systems*, vol. 34, no. 8, pp. 3779–3795, 2023.
- [9] H. Van Hasselt, A. Guez, and D. Silver, “Deep Reinforcement Learning with Double Q-learning,” *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 30, no. 1, 2016.